



Introduction to Machine Learning

CSCE 478/878

Programming Assignment 4

Fall 2020

Linear Support Vector Machine & Principle Component Analysis

Basic Info

You will work in teams of maximum three students from the previous assignment.

The programming code will be graded on **both implementation and correctness**.

This assignment **doesn't require a written report**.

Assignment Goals

This assignment is intended to build the following skill:

- Implement the Gradient Descent algorithm for the Linear Support Vector Machine classifier model.
 - Perform Principle Component Analysis (PCA) based dimensionality reduction by using the eigendecomposition technique.
-

Assignment Instructions

Note: you are not allowed to use any Scikit-Learn or python library for building the Linear SVM model and performing PCA.

- The code should be written in a Jupyter notebook. Use the following naming convention.
`<lastname1>_<lastname2>_<lastname3>_assignment4.ipynb`
 - The Jupyter notebook should be submitted via webhandin.
-

Score Distribution

Part A: 478 (70 pts) & 878 (80 pts)

Part B: 478 (45 pts) & 878 (45 pts)

Total: 478 (115 pts) & 878 (125 pts)

Part A: Linear Support Vector Machine (SVM)

Dataset: You will use the Iris dataset for **binary classification**. Use petal length and petal width features of the Iris dataset, and determine whether a sample is Iris Virginica or not.

URL: https://scikit-learn.org/stable/modules/generated/sklearn.datasets.load_iris.html

1. Implement a **Linear_SVC** model class for performing **binary classification**. The model should implement the batch Gradient Descent (GD) algorithm.

[40 pts]

a) `__init__(self, C=1, max_iter=100, tol=None, learning_rate='constant', learning_rate_init=0.001, t_0=1, t_1=1000, early_stopping=False, validation_fraction=0.1, **kwargs)`

This method is used to initialize the data members of the class when an object of class is created. For example, `self.C = C`

Arguments:

`C : float`

It provides the regularization/penalty coefficient.

`max_iter : int`

Maximum number of iterations. The GD algorithm iterates until convergence (determined by 'tol') or this number of iterations.

`tol : float`

Tolerance for the optimization.

`learning_rate : string (default 'constant')`

It allows to specify the technique to set the learning rate: constant learning for all iterations, or varying learning rate given by a learning rate schedule function.

- 'constant': a constant learning rate given by 'learning_rate_init'.
- 'adaptive': gradually decreases the learning rate based on a schedule. Write a function that would be used if learning_rate is set to 'adaptive'. When 'adaptive' is used, the 'learning_rate_init' parameter has no effect as the learning rate varies by a learning rate schedule function. It uses the 't_0' and 't_1' parameters (see below).

Pseudocode for the “adaptive” learning_rate function:

Write a function that decreases learning rate gradually during each iteration:

$$\text{Learning rate} = \frac{t_0}{\text{iteration} + t_1}$$

where t_0 and t_1 are two constants that you need to determine empirically.

Choose constant t_0 and t_1 such that initially the learning rate is large enough.

learning_rate_init : double

The initial learning rate value if learning_rate is set to 'constant'. It controls the step-size in updating the weights. It has no effect if the 'learning_rate' is 'adaptive'.

early_stopping : Boolean, default=False

Whether to use early stopping to terminate training when validation score is not improving. If set to True, it will automatically set aside a fraction of training data as validation and terminate training when validation score is not improving.

validation_fraction : float, default=0.1

The proportion of training data to set aside as validation set for early stopping. Must be between 0 and 1. Only used if early_stopping is True.

b) fit(self, X, Y):

Implement the batch GD algorithm in the fit method. The weight vector and the intercept/bias should be denoted by w and b, respectively. Store the cost values for each iteration so that later you can use it to create a learning curve.

Arguments:

X : ndarray

A numpy array with rows representing data samples and columns representing features.

Y : ndarray

A 1D numpy array with labels corresponding to each row of the feature matrix X.

Note: the “fit” method should update the following parameters:

```
self.intercept_ = np.array([b])
self.coef_ = np.array([w])
self.support_vectors_ =
```

The “fit” method should display the total number of iterations using a print statement.

Returns:
self

c)
predict(self, X)

Arguments:
X : ndarray
A numpy array containing samples to be used for prediction. Its rows represent data samples and columns represent features.

Returns:
1D array of predicted class labels for each row in X.

Note: the “predict” method uses the **self.coef_[0]** and **self.intercept_[0]** to make predictions.

Binary Classification using Linear_SVC Classifier

2. Read the Iris data using the [sklearn.datasets.load_iris](#) method. Create the data matrix X by using two features: petal length and petal width. Recode the binary target such that Iris-Virginica samples are 1, and other samples are 0.
[1 pts]
3. Partition the data into train and test set (80% - 20%). Use the “**Partition**” function from your previous assignment.
[2 pts]
4. **Model selection via Hyper-parameter tuning:** Use the **kFold** function from previous assignment to find the optimal values for the following hyperparameters.
[5 pts]

C
learning_rate
learning_rate_init (when ‘constant’ learning_rate is used)

`max_iter`
`tol`

5. Train the model using optimal values for the hyperparameters and evaluate on the **test data**. Report the test accuracy and test confusion matrix.
[5 pts]
6. Plot the learning curve.
[5 pts]
7. Plot the decision boundary and show the support vectors using the “`decision_boundary_support_vectors`” function given in:
<https://github.com/rhasanbd/Support-Vector-Machine-Classifiers-Beginners-Survival-Kit/blob/master/Support%20Vector%20Machine-1-Linearly%20Separable%20Data.ipynb>
[12 pts]

Note that if your test accuracy is less than 95% you will lose 10% of the total obtained points. If your test accuracy is less than 90% you will lose 30% of the total obtained points.

8. [**Extra Credit for 478 and Mandatory for 878**] Implement early stopping in the “fit” method of the Linear_SVC model. You will have to use the following two parameters of the model: `early_stopping` and `validation_fraction`. Also note that when training the model using early stopping it should generate an early stopping curve.
[10 pts]

Part B: Principle Component Analysis

You will perform dimensionality reduction on a grayscale image (posted on Canvas) using PCA. The PCA will be implemented using the **eigendecomposition** technique. More specifically, you will find the top k eigenvectors (i.e., principle components) of a pixel matrix (a gray scale image). Then, using the top k eigenvector matrix, you will project the pixel matrix on its principle components. This will reduce the dimension of the pixel matrix without losing much variance.

Note: See the following notebook for understanding the manual implementation of the eigendecomposition based PCA using python.
<https://github.com/rhasanbd/Dimensionality-Reduction-Get-More-From-Less-And-See-the-Unseen/blob/master/Dimensionality%20Reduction-PCA-Eigendecomposition-Introduction.ipynb>

9. Using the matplotlib.pyplot “*imread*” function read the image as a 2D matrix. Denote it with “X”. Show the image using matplotlib.pyplot *imshow* function.

If the image is RGB, then you need to convert it into a grayscale image, as follows (use matplotlib.pyplot “gray” function).

```
X = imread("image_path")[:, :, 0]
gray()
```

[3 pts]

10. Implement the steps of eigendecomposition based PCA on X: (a) mean center the data matrix X, (b) compute the covariance matrix from it, (c) find eigenvalues and eigenvectors of the covariance matrix (you may use the *numpy.linalg.eig* function). **[7 pts]**

11. Then, find the top k eigenvectors (sort eigenvalue-eigenvector pairs from high to low, and get the top k eigenvectors), and create an eigenvector matrix using top k eigenvectors (each eigenvector should be a column vector in the matrix, so there should be k columns). **[10 pts]**

12. Finally project the mean centered data on the k top eigenvectors (it should be a dot product between mean centered X and the top k eigenvector matrix).

[5 pts]

13. Reconstruct the data matrix by taking dot product between the projected data (from last step) and the transpose of the top k eigenvector matrix.

[5 pts]

14. Compute the reconstruction error between the mean centered data matrix X and reconstructed data matrix (you may use the *sklearn.metrics.mean_squared_error* function).

[5 pts]

15. Perform steps 11 – 14 for the following values of k : 10, 30, 50, 100, 500. For each k , show the reconstructed image (use the matplotlib.pyplot *imshow* function with the reconstructed data matrix for each k). With each reconstructed image print the value of k and the reconstruction error.

[10 pts]